Sayon Sivakumaran

Email: sayon.sivakumaran@mail.utoronto.ca
Website: sayonsivakumaran.github.io
LinkedIn: linkedin.com/in/sayonsivakumaran
GitHub: github.com/sayonsivakumaran

EDUCATION

University of Toronto | Bachelor of Science, Computer Science & Statistics **Sep. 2018 – Apr. 2023**

- Specialist Programs in Software Engineering & Data Science, **cGPA**: 3.79/4.00
- **Relevant Courses**: Software Engineering, Data Structures & Algorithms, Systems Programming
- **Teaching Assistant**: Introduction to Probability

EXPERIENCE

Amazon | Software Development Engineer Intern **Jul. 2021 – Oct. 2021**

- Deprecated an Amazon EC2 fleet that was only used for handling large requests by re-directing those requests towards a central fleet, saving approximately $1,000 yearly in fleet expenses
- Leveraged an internal tool to throttle incoming requests to a service based on the available capacity of a host, while also adding priority to accept small requests, decreasing the chance of a host having full capacity by over 50%

Okta | Software Engineer Intern **May 2021 – Jul. 2021**

- Implemented a UI using Java, JavaScript & MySQL, allowing App Analysts to update an application's metadata in production directly through the interface, decreasing average deployment time for metadata changes by 5+ days
- Enhanced an existing API endpoint to allow specific users to post changes to an application's classification categories

Georgian | Machine Learning Engineer Intern **Jan. 2021 – Apr. 2021**

- Extended cloud catalog to support using AWS Batch within hydra – a machine learning tool used to execute & optimize experiments using different cloud platforms – decreasing average setup time for an individual experiment by 5+ minutes
- Designed infrastructure-as-code to host an auto-scaling MLflow tracking server on AWS & GCP using Terraform

Bond Brand Loyalty | Data Engineer Intern **Sep. 2020 – Dec. 2020**

- Implemented transformation logic for 20+ SQL tables in ELT pipelines from the on-site database to Azure Blob Storage using Python & Spark
- Optimized the Internal Finance Operations ELT pipeline in Microsoft Azure so that all data transformations are performed within one Databricks cluster, reducing daily total cluster start-up time by over 90%

Leonardo Worldwide Corporation | Full-Stack Engineer Intern **Aug. 2019 – Apr. 2020**

- Kicked off the development of a data migration tool using Spring Boot & MySQL to migrate information for 1,000+ customers & 100+ hotels from a legacy system to a modern platform
- Developed an API endpoint used to grant admin accounts control over receiving automated emails based on specific user events using Java, gRPC & GraphQL, decreasing sent emails by over 40%

TECHNICAL SKILLS

Languages: Java, Python, JavaScript, SQL, C/C++, HTML, Kotlin, R

Frameworks: Spring Boot, React.js, Spark, Hibernate, gRPC, GraphQL, Kafka, Express.js, Selenium

Other: Terraform, AWS, GCP, Microsoft Azure, MongoDB, Docker, Databricks, Snowflake, NGINX

PROJECTS

Scarborough Dining – Website that acts as a central platform to share information about local restaurants

- Selected by City of Toronto as foundational code for the Find Dining program – a community initiative to provide an online marketing and networking platform for local restaurant owners in East Toronto

University Valuation Research Report – Statistical report that evaluates if a university education is worth its cost

- Analyzed social and financial factors by rearranging and plotting data to create various statistical models